



 @AndrzejWasowski

Andrzej Wąsowski
Florian Biermann

Advanced Programming

Purely Functional Parallelism [Par]

IT UNIVERSITY OF COPENHAGEN

S SOFTWARE
Q QUALITY
R RESEARCH

Study lab:

DR 2F14

4A58: live exercise solving (14:30)

3A52: exercise session as usual

Parallel Data Structures (Motivation)

```
1 val list = (1 to 100).toList
2 val flist = list.map(f)
```

- If f slow but referentially transparent, parallelizing is beneficial!
- The parallel collections library for Scala allows: `list.par.map (f)`
- `list.par: collection.parallel.immutable.ParSeq[Int] = ParVector(1, ...`
- More types supported: `ParArray`, `ParVector`, `mutable.ParHashMap`, `mutable.ParHashSet`, `immutable.ParHashMap`, `immutable.ParHashSet`, `ParRange`, `ParTrieMap`
- **Similarities** with `Par`: enable parallelism at the level of processing data structures without low level concurrency primitives (**parallel programming for the masses!**)
- **Differences** from `Par`: Scala's parallel collections are **eager**. We separate construction of the computation from execution. This gives more flexibility.
- Similar facilities exist in LINQ (C#) and in F#

Parallel Collections in Spark

- Spark has similar facilities:

```
val data = Array(1, 2, 3, 4, 5)
```

```
val distData = sc.parallelize(data)
```

- Constructs an RDD from a collection.
- RDD resembles a parallel collection, but it can **also** be distributed
- **RDD constructions are lazy.** As long as transformations are applied to an RDD, no computation is executed.
- Allows Spark schedulers to control the computation better
- This is more like `Par` than Scala's native parallel collections
- Today we look at a design of such general APIs

API for Functional Parallelism

What problem we are solving today?

The Problem:

- Design a general functional parallel programming API
- That allows to arbitrarily structure a parallel computation
- And then to execute it on a limited amount of resources (a bounded thread pool)

The Solution:

- **Separate** description of computation from execution
- Use **Continuation Passing Style** (CPS) to pass resources between threads without blocking
- Use of familiar combinators to **combine computations** (flatMap, map, map2, sequence, ...). The semantics of these combinators is much less trivial than for list, option, or state monads.

API for Functional Parallelism

How to read this chapter?

Chapter 7

- No right answers in design
- You will see a collection of design choices
- You are to understand their trade-offs, and think critically.
- The lecture explains the key design, but not all the meanders of the story — not the full learning experience!

Agenda for Today

- Motivation for Par [Done]
- Usage & Design of Par
- Continuation Passing Style (A general pattern)
- Implementation of Par
- Extension methods / Opaque types

map2 for Option and Par

```
def map2[A, B, C](oa: Option[A], ob: => Option[B])(f: (A, B) => C): Option[C] =  
  oa.flatMap { a => ob.map { b => f(a, b) } }
```

- Why is oa by-value and ob by-name?
- Now a version for Par:

```
def map2[A, B, C](pa: => Par[A], pb: => Par[B])(f: (A, B) => C): Par[C]
```

- Why are both pa and pb by name?
- An example of use, parallel summation of list:

```
1 def sum(ints: IndexedSeq[Int]): Par[Int] =                               // Returns immediately  
2   if ints.size <= 1 then  
3     Par.unit(ints.headOption.getOrElse(0))  
4   else  
5     val (l, r) = ints.splitAt(ints.length/2)  
6     Par.map2(Par.fork(sum(l)), Par.fork(sum(r)))(_ + _)                // After forking don't wait explicitly!  
7                                                                           // We can apply _+_ to computations!
```


The Par Type

- `Par[A]`: a **pure data structure** describing a **parallel computation**
 - Some similarity to `Stream` which describes computations happening in sequence
 - `Par` is Java's `Callable` with a way better API
 - Allows expanding the computation, without waiting for results
 - Separates construction and parallelization of computation from scheduling and execution
 - First decide what runs in separate threads, then on what resource (`Executor`) to run it
- `unit[A](a: A): Par[A]` promote a constant to `Par` eagerly (trivial, return a immediately)
- `map2[A, B, C](pa: Par[A], pb: Par[B])(f: (A, B) => C): Par[C]` combines results of two computations with a binary function. Does not introduce new threads.
- `fork[A](a: => Par[A]): Par[A]` marks a computation for concurrent evaluation (separate thread). No evaluation until forced by `run`. Executes on a different thread.
- `lazyUnit[A](a: => A): Par[A]` wraps its unevaluated argument in a `Par` and marks it for concurrent evaluation (so it combines `unit` and `fork`).
- `run[A](es: ExecutorService)(p: Par[A]): A` extracts a value from a `Par` by actually performing the computation (the non-blocking version)

Other Combinators

- `def map[A, B](pa: Par[A])(f: A => B): Par[B] = map2(pa, unit(())) { (a, _) => f(a) }`
 - `map` **extends** the definition of a parallel computation with a **new step** (`f`)
 - Example: `def sortPar(parList: Par[List[Int]]) = map(parList) { _.sorted }`
 - Q. What does `sortPar` do?
 - It changes a parallel computation producing a list, into one whose resulting list is sorted.
Nothing is run at this stage!
 - An example how we can build a computation without committing to where and how to execute
 - Choose to use a different number of threads or a different scheduling policy in different places
 - For example UI vs background batch processing
- `def asyncF[A, B](f: A => B): A => Par[B]` change `f` to **run in parallel** (exercise)
- `def sequence[A](ps: List[Par[A]]): Par[List[A]]`
 - **Recomposes** a list of parallel computations of `A` into a single parallel computation of a list
 - Example: schedule `n` downloads, get a list of downloaded files in parallel (exercise)
 - Familiar from `State`
- We can use it to implement `parMap` (that maps over list in parallel):
`def parMap[A, B](as: List[A])(f: A => B): Par[List[B]] =
 sequence(as.map(asyncF(f)))`

Mentimeter

```
■ def map[A, B](pa: Par[A])(f: A => B): Par[B] = map2(pa, unit(())) (a, _) => f(a)
■ def asyncF[A, B](f: A => B): A => Par[B]
■ def sequence[A](ps: List[Par[A]]) : Par[List[A]]
■ def parMap[A, B](as: List[A])(f: A => B): Par[List[B]] =
    sequence(as.map(asyncF(f)))
```

Question. Why is `asyncF` called? What would happen if we did:

```
def parMap[A, B](as: List[A])(f: A => B): Par[List[B]] =
    sequence(as.map(f))
```

Continuation Passing Style

```
1  def f(x: X): Y                                // Consider two functions
2  def g(y: Y): Z
3  // Normally this is how we compose them:
4  g(f(x)): Z                                    // f (x) first returns, then we call 'g' on the outcome

6  def f(x: X)(cont: Y => Z): Z =                 // Rewrite f to call a *continuation* instead of returning
7      val y_result = ...                        // the original body of f
8      cont(y_result)                           // tail call, no new stack/thread use


10 def g[A](y: Y)(cont: Z => A): A =              // Do the same to 'g';
11     val z_result = ...                        // the original body of g
12     cont(z_result)

14 // Let 'consumer: Z => Unit' execute what we shall do with g's result
15 f(x) { y => g(y)(consumer) }                  // Compose 'f' and 'g' in the continuation passing style
```

- All returning via argument passing; Unit inessential, consumer could be pure and return value
- The last thing each function does is calling the next function
- A peculiar generalization of tail recursion and accumulator style
- Used to implement Par so that handing over from one to another can reuse same thread

Background: Java's ExecutorService and Future API

```
class ExecutorService {  
    def submit[A](a: Callable[A]): Future[A]  
}  
  
trait Future[A] {  
    def get: A  
}
```




https://www.slideshare.net/afkham_azeez/java-colombo-developing-highly-scalable-apps

- An ExecutorService provides methods to **manage termination**, and
- Can produce a Future for **tracking progress of asynchronous tasks**
- A **thread pool** is a **blocking queue**
- A **worker thread** executes a task from the queue as long as non-empty

Typical use of an ExecutorService

```
class ExecutorService {  
    def submit[A](a: Callable[A]): Future[A]  
}  
  
trait Future[A] {  
    def get: A  
}
```



Typical usage1

```
val service: ExecutorService =  
    Executors.newFixedThreadPool(2)  
service.submit(t1)  
service.submit(t2)  
service.submit(t3)
```

The Implementation of Par [1/3]

The non-blocking CPS-based variant (Section 7.4.4)

```
1 // We use futures to represent an asynchronous calculation of a value.
2 // Java Futures don't have a way to continue computation without waiting for A. Par has.
3 opaque type Future[+A] = (A => Unit) => Unit           // Java futures deadlock, new design
4                                                       // The future calls the continuation when ready

6 opaque type Par[A] = ExecutorService => Future[A]    // Just an alias, using Java's Executor
7 // opaque means it looks like a class outside this file (others cannot see it is a function)

9 // Normally we do not execute Par, but compose it with new calculations (using map,map2,chooser,etc.)
10 // Once we have a representation of the whole thing we can run it:
11 extension [A] (pa: Par[A])                          // An extension method for Par[A] (only)
12   def run[A](es: ExecutorService): A =
13     val ref = java.util.concurrent.atomic.AtomicReference[A]() // Mutable threadsafe cell (local!)
14     val latch = CountDownLatch(1)                             // Create a latch.
15     pa (es) { a => ref.set(a); latch.countDown }              // Continuation sets ref and
16                                                                // decrements latch.
17     latch.await                                                // Wait for latch reaching 0
18                                                                // (never if p crashes.)
19     ref.get                                                     // Return 'a' set by the continuation.
```

The Implementation of Par [2/3]

The non-blocking CPS-based variant (Section 7.3.4)

```
1 def unit[A](a: A): Par[A] =                               // A strict unit.
2   es => k => k (a)

4 def eval(es: ExecutorService)(r: => Unit): Unit =         // A helper function.
5   es.submit(new Callable[Unit] { def call = r })          // Submit a unit computation to an executor.

7 def fork[A](a: => Par[A]): Par[A] =                       // Marks 'a' for parallel execution.
8   es => k => eval(es)(a(es)(k))                          // Do not evaluate, but delay in a Future, and eval.

10 def lazyUnit[A](a: => A): Par[A] =                       // A lazy (by-name) version of unit.
11   fork(unit(a))                                           // Mark 'a' for parallel, wrap in unit.
12                                                         // NB: fork is by-name.
```


The Implementation of Par [3/3]

The non-blocking CPS-based variant (Section 7.4.4)

```
1 extension [A] (pa: Par[A])
2   def map2[B, C](pb: Par[B])(f: (A, B) => C): Par[C] =
3     es => k =>                                     // Common to use 'k' for "continuation".
4       var ar: Option[A] = None                       // NB: var!
5       var br: Option[B] = None                       // NB: var!
6       val combiner = Actor[Either[A, B]] (es) {      // Shorthand for x => x match ...
7         case Left (a) =>
8           if br.isDefined then Par.eval(es)(k(f(a,br.get))) // Continue with result from b.
9           else ar = Some(a)                               // No result for b, provide result for a.
10        case Right(b) =>
11          if ar.isDefined then Par.eval(es)(k(f(ar.get,b))) // Continue with result from a.
12          else br = Some(b)                               // No result for a, provide result for b.
13      }
14      pa(es) { a => combiner ! Left (a) }              // Returns Unit;
15      pb(es) { b => combiner ! Right (b) }            // the real result is the function!
```

**Other operators are derived
(see exercises)**

Extension Methods (C# vs Scala)

```
1 namespace ExtensionMethods {  
2     public static class MyExtensions {  
3         public static int WordCount(this String str)  
4             => str  
5                 .Split(  
6                     new[] { ' ', '.', '?' },  
7                     StringSplitOptions.RemoveEmptyEntries)  
8                 .Length;  
9     }  
10 }  
11 using ExtensionMethods;  
12 "Hello Extension Methods".WordCount();
```

```
1 extension (val str: String)  
2     def wordCount =  
3         str.split(" .?".toArray)  
4             .filter { !_.isEmpty }  
5             .length  
  
11 import wordCount  
12 "Hello Extension Methods".wordCount
```

- **Extension methods** C#, F#, Xtend, Kotlin: define static methods, call like instance method
- That's why String in Scala has **more methods** than in Java, even though it is the **same class**!
- In fact, split is a method on StringOps not on String (see above)
- **Extensions work very well with opaque types**: limit an extension only to a named type

Extension Methods

- A mechanism to **extend an existing library**
- When you **cannot change the source** code
- Add methods to classes **without recompiling** the source
- **Even to Java classes from 1995!**
- Add methods to classes at **call location**, not at class definition location
- Even **objects** produced by **old code** (factories) get the new methods
- When you **read** someone else's code you need to know that you have to search **not only for class methods** but also for extension methods
- In Scala, extensions are often placed in the *Ops classes, e.g. <https://www.scala-lang.org/api/2.12.3/scala/collection/immutable/StringOps.html>
- Warning: in older versions of Scala, implicits have been used to implement extensions. A more complex syntax and mechanism to explain. These are now deprecated. Beware Stack Overflow!
- **Exercises:** use this pattern to add methods to Par which is a function type alias! Not even explicitly a class!